



```
In [1]: ▶ import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.simplefilter("ignore")
```

```
In [2]: ▶ df = pd.read_csv('gene_expression.csv')
df.head()
```

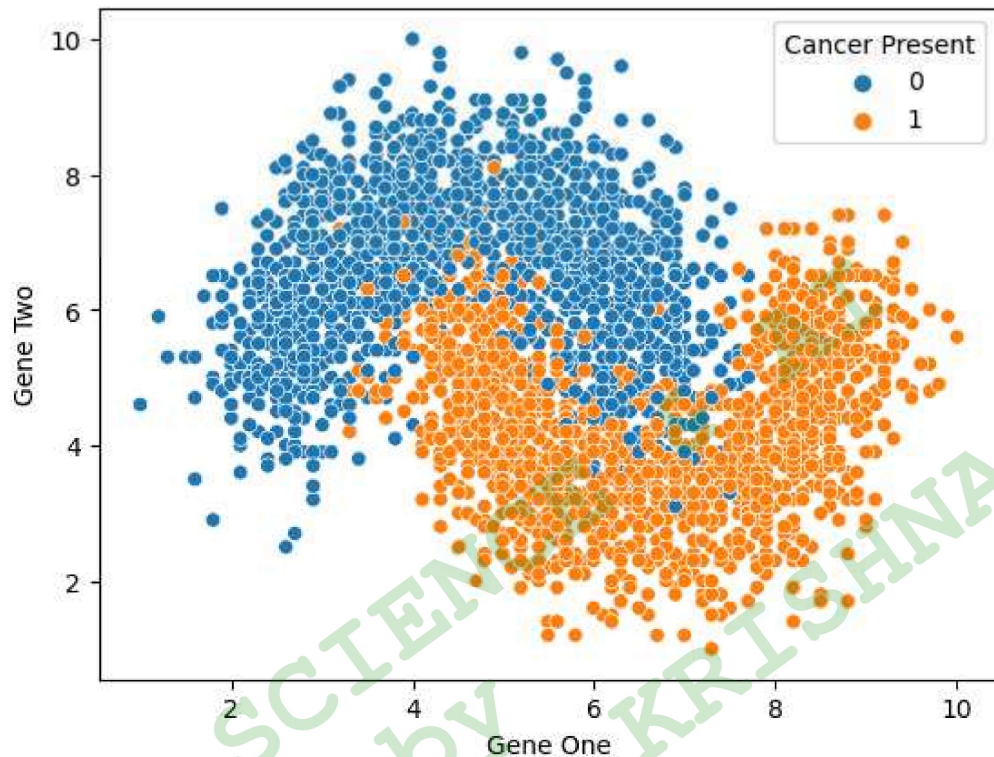
Out[2]:

	Gene One	Gene Two	Cancer Present
0	4.3	3.9	1
1	2.5	6.3	0
2	5.7	3.9	1
3	6.1	6.2	0
4	7.4	3.4	1

```
In [3]: ▶ df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Gene One        3000 non-null   float64
1   Gene Two        3000 non-null   float64
2   Cancer Present  3000 non-null   int64
dtypes: float64(2), int64(1)
memory usage: 70.4 KB
```

```
In [4]: sns.scatterplot(x='Gene One',y='Gene Two',hue='Cancer Present',data=df)
plt.show()
```



```
In [5]: df.isnull().sum()
```

```
Out[5]: Gene One      0
Gene Two      0
Cancer Present  0
dtype: int64
```

**X & y**

```
In [6]: X = df.drop('Cancer Present',axis=1)
y = df['Cancer Present']
```

**Train|Test Split**

```
In [7]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size=0.2,
                                                    random_state=9)
```

**Scaling Data**



```
In [8]: ▶ from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()

        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)
```

### Hyperparameter Tuning for KNN Classifier

```
In [9]: ▶ from sklearn.neighbors import KNeighborsClassifier
        estimator = KNeighborsClassifier()
        param_grid = {'n_neighbors': list(range(1,100))}

        from sklearn.model_selection import GridSearchCV
        cv_classifier = GridSearchCV(estimator, param_grid,
                                     cv=5, scoring='accuracy')

        cv_classifier.fit(X_train, y_train)

        cv_classifier.best_params_
```

```
Out[9]: {'n_neighbors': 19}
```

### KNN model with best hyper parameters

```
In [10]: ▶ #Modelling
        from sklearn.neighbors import KNeighborsClassifier
        knn = KNeighborsClassifier(n_neighbors=19)
        knn.fit(X_train, y_train)

        # Prediction
        y_pred_test = knn.predict(X_test)
        y_pred_train = knn.predict(X_train)

        #Evaluation
        from sklearn.metrics import accuracy_score
        print("Train accuracy:", accuracy_score(y_train, y_pred_train))
        print("Test accuracy:", accuracy_score(y_test, y_pred_test))

        from sklearn.model_selection import cross_val_score
        print("Cross Validation Score:", cross_val_score(knn, X, y, cv=5).mean())
```

```
Train accuracy: 0.9408333333333333
Test accuracy: 0.9333333333333333
Cross Validation Score: 0.9313333333333335
```

```
In [11]: ▶ from sklearn.metrics import confusion_matrix
        confusion_matrix(y_test, y_pred_test)
```

```
Out[11]: array([[264, 19],
                [ 21, 296]], dtype=int64)
```



```
In [12]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred_test))
```

	precision	recall	f1-score	support
0	0.93	0.93	0.93	283
1	0.94	0.93	0.94	317
accuracy			0.93	600
macro avg	0.93	0.93	0.93	600
weighted avg	0.93	0.93	0.93	600

## Prediction on New data

### New Data

```
In [13]: df = pd.DataFrame({"Gene One": [4.9], "Gene Two": [3.9]})
df
```

Out[13]:

	Gene One	Gene Two
0	4.9	3.9

### preprocess the new data

```
In [14]: df_scaled = scaler.transform(df)
```

### use KNN model to predict on new data

```
In [15]: knn.predict(df_scaled)
```

Out[15]: array([1], dtype=int64)